

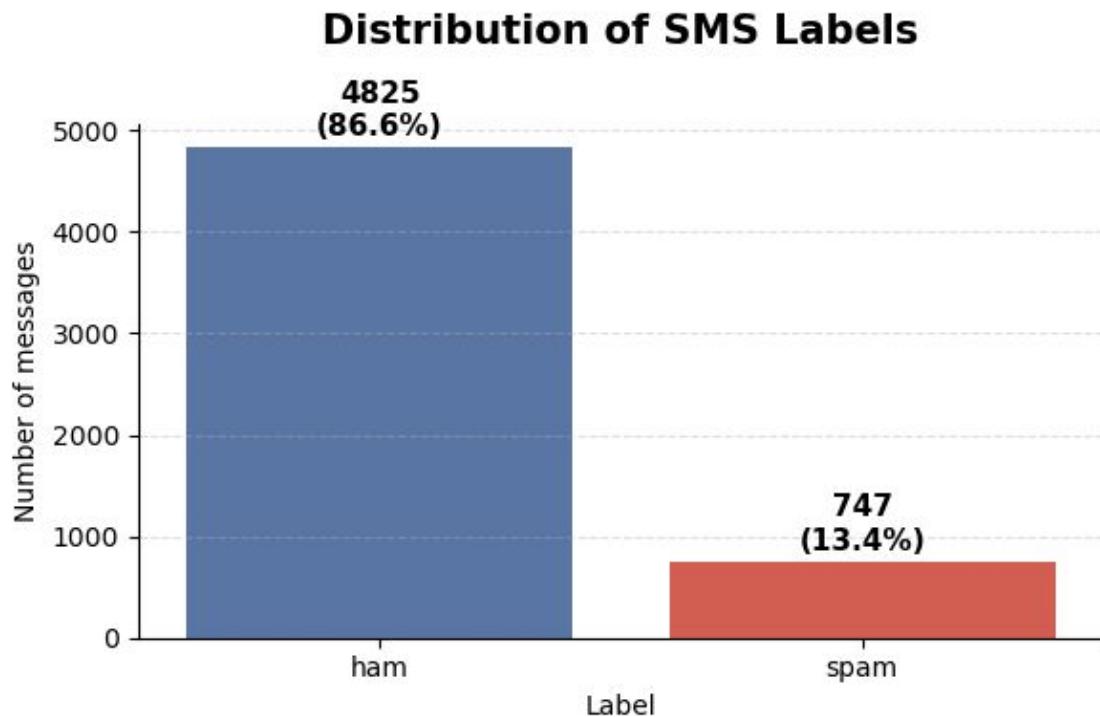


Spam Detector

Automatically classify SMS messages as spam or ham

Project Overview

5572 messages



Data Cleaning & Preprocessing

- ❖ **Cleaning:**
 - lower casing
 - removing punctuation
 - keeping digits
- ❖ **Tokenization:**
 - split text into words,
 - padded/truncated to 40 tokens
- ❖ **Vocabulary:**
 - 7 600 unique words

Model Architecture


- *Embedding(128)* → learns semantic meaning of words
- *LSTM(64)* → captures the sequence of words
- *Dropout(0.3)* → reduces overfitting
- *Dense(1, sigmoid)* → outputs spam probability

Layer (type)	Output Shape
text_vectorization_4 (TextVectorization)	(1, 40)
embedding (Embedding)	?
lstm_4 (LSTM)	?
dropout_4 (Dropout)	?
dense_4 (Dense)	?

Training Results: heatmap

Each cell shows how well the model performs for **ham** and **spam** in terms of **precision**, **recall**, and **F1-score**.

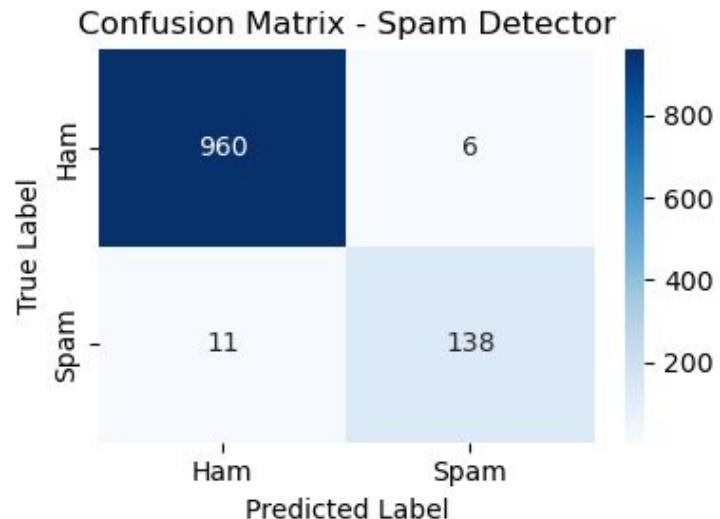
Classification Report Summary



F1-score combines **precision** and **recall** – near-perfect values indicate a robust model.

Training Results: Confusion matrix

- Each cell shows the number of messages falling into a specific category of prediction vs reality.



True Negative Real <i>ham</i> predicted as <i>ham</i>	False Positive Real <i>ham</i> predicted as <i>spam</i>
False Negative Real <i>spam</i> predicted as <i>ham</i>	True Positive Real <i>spam</i> predicted as <i>spam</i>

Error Analysis Examples

False Positives

(predicted spam but actually ham):

- [0.87] waiting for your **call**
- [0.81] **nokia phone** is lovly
- [0.81] height of confidence all the aeronautics professors wer **calld** amp they wer askd 2 sit in an aeroplane aftr they sat they...
- [0.52] **unlimited texts** limited minutes

False Negatives

(predicted ham but actually spam):

- [0.00] sorry i missed your **call** let's talk when you have the time i'm on **07090201529**
- [0.00] for **sale** arsenal dartboard good condition but no doubles or trebles
- [0.18] latest **news** police station toilet stolen cops have nothing to go on

True Negative Real <i>ham</i> predicted as <i>ham</i>	False Positive Real <i>ham</i> predicted as <i>spam</i>
False Negative Real <i>spam</i> predicted as <i>ham</i>	True Positive Real <i>spam</i> predicted as <i>spam</i>

Prediction Tool: Real-time Spam Detection

This simple prediction interface allows testing the model with new unseen messages.

- The tool takes raw text input (SMS) and predicts its **probability of being spam**.
- Messages are color-coded: **green for HAM, red for SPAM**.

	Message	Predicted Label	Spam Probability
0	Congratulations! You've won a new iPhone, click here to claim!	SPAM	94.97%
1	Hi John, can you send me the report by tomorrow?	HAM	0.06%
2	Urgent! Your bank account has been locked, verify immediately.	SPAM	90.82%
3	Ok cool, I'll bring the cake for Saturday.	HAM	0.21%



Thank you for your attention
– any questions?

